

On-device inference --- model fetch, checksum-verified cache, and prediction

